

TLA+ specification of *Lamport's* distributed mutual-exclusion algorithm that appeared as an example in L. *Lamport*: Time, Clocks and the Ordering of Events in a Distributed System. *CACM* 21(7):558 – 565, 1978.

EXTENDS *Naturals*, *Sequences*

The parameter  $N$  represents the number of processes. The parameter  $maxClock$  is used only for model checking in order to keep the state space finite.

CONSTANT  $N$ ,  $maxClock$

ASSUME  $NType \triangleq N \in Nat$

ASSUME  $maxClockType \triangleq maxClock \in Nat$

$Proc \triangleq 1 \dots N$

$Clock \triangleq Nat \setminus \{0\}$

For model checking, add *ClockConstraint* as a state constraint to ensure a finite state space and override the definition of *Clock* by  $1 \dots maxClock + 1$  so that *TLC* can evaluate the definition of *Message*.

VARIABLES

$clock$ ,	local clock of each process
$req$ ,	requests received from processes (clock transmitted with request)
$ack$ ,	acknowledgements received from processes
$network$ ,	messages sent but not yet received
$crit$	set of processes in critical section

Messages used in the algorithm.

$ReqMessage(c) \triangleq [type \mapsto \text{"req"}, clock \mapsto c]$

$AckMessage \triangleq [type \mapsto \text{"ack"}, clock \mapsto 0]$

$RelMessage \triangleq [type \mapsto \text{"rel"}, clock \mapsto 0]$

$Message \triangleq \{AckMessage, RelMessage\} \cup \{ReqMessage(c) : c \in Clock\}$

The type correctness predicate.

$TypeOK \triangleq$

$clock[p]$  is the local clock of process  $p$   
 $\wedge clock \in [Proc \rightarrow Clock]$   
 $req[p][q]$  stores the clock associated with request from  $q$  received by  $p$ , 0 if none  
 $\wedge req \in [Proc \rightarrow [Proc \rightarrow Nat]]$   
 $ack[p]$  stores the processes that have *ack*'ed  $p$ 's request  
 $\wedge ack \in [Proc \rightarrow SUBSET Proc]$   
 $network[p][q]$ : queue of messages from  $p$  to  $q$  – pairwise *FIFO* communication  
 $\wedge network \in [Proc \rightarrow [Proc \rightarrow Seq(Message)]]$   
 $crit$  set of processes in critical section: should be empty or singleton  
 $\wedge crit \in SUBSET Proc$

$$Init \triangleq$$

*beats*( $p, q$ ) is true if process  $p$  believes that its request has higher priority than  $q$ 's request. This is true if either  $p$  has not received a request from  $q$  or  $p$ 's request has a smaller clock value than  $q$ 's. If there is a tie, the numerical process  $ID$  decides.

Broadcast a message: send it to all processes except the sender.

Process  $p$  requests access to critical section.

Process  $p$  receives a request from  $q$  and acknowledges it.

Process  $p$  receives an acknowledgement from  $q$ .

2

$$\begin{aligned}
& \wedge ack' = [ack \text{ EXCEPT } ![p] = @ \cup \{q\}] \\
& \wedge network' = [network \text{ EXCEPT } ![q][p] = Tail(@)] \\
& \wedge \text{UNCHANGED } \langle clock, req, crit \rangle
\end{aligned}$$

Process  $p$  enters the critical section.

$$\begin{aligned}
Enter(p) & \triangleq \\
& \wedge ack[p] = Proc \\
& \wedge \forall q \in Proc \setminus \{p\} : beats(p, q) \\
& \wedge crit' = crit \cup \{p\} \\
& \wedge \text{UNCHANGED } \langle clock, req, ack, network \rangle
\end{aligned}$$

Process  $p$  exits the critical section and notifies other processes.

$$\begin{aligned}
Exit(p) & \triangleq \\
& \wedge p \in crit \\
& \wedge crit' = crit \setminus \{p\} \\
& \wedge network' = [network \text{ EXCEPT } ![p] = Broadcast(p, RelMessage)] \\
& \wedge req' = [req \text{ EXCEPT } ![p][p] = 0] \\
& \wedge ack' = [ack \text{ EXCEPT } ![p] = \{\}] \\
& \wedge \text{UNCHANGED } clock
\end{aligned}$$

Process  $p$  receives a release notification from  $q$ .

$$\begin{aligned}
ReceiveRelease(p, q) & \triangleq \\
& \wedge network[q][p] \neq \langle \rangle \\
& \wedge \text{LET } m \triangleq Head(network[q][p]) \\
& \quad \text{IN } \wedge m.type = \text{"rel"} \\
& \quad \wedge req' = [req \text{ EXCEPT } ![p][q] = 0] \\
& \quad \wedge network' = [network \text{ EXCEPT } ![q][p] = Tail(@)] \\
& \quad \wedge \text{UNCHANGED } \langle clock, ack, crit \rangle
\end{aligned}$$

Next-state relation.

$$\begin{aligned}
Next & \triangleq \\
& \vee \exists p \in Proc : Request(p) \vee Enter(p) \vee Exit(p) \\
& \vee \exists p \in Proc : \exists q \in Proc \setminus \{p\} : \\
& \quad ReceiveRequest(p, q) \vee ReceiveAck(p, q) \vee ReceiveRelease(p, q) \\
vars & \triangleq \langle req, network, clock, ack, crit \rangle \\
Spec & \triangleq Init \wedge \Box [Next]_{vars}
\end{aligned}$$

A state constraint that is useful for validating the specification using finite-state model checking.

$$ClockConstraint \triangleq \forall p \in Proc : clock[p] \leq maxClock$$

No channel ever contains more than three messages. In fact, no channel ever contains more than one message of the same type, as proved below.

$BoundedNetwork \triangleq \forall p, q \in Proc : Len(network[p][q]) \leq 3$

The main safety property of mutual exclusion.

$Mutex \triangleq \forall p, q \in crit : p = q$

$Live \triangleq \forall p \in Proc : \Diamond(p \in crit)$

---

\ \* Modification History

\ \* Last modified Sun Oct 07 00:34:02 EDT 2018 by mehuljain

\ \* Created Sun Sep 30 20:38:42 EDT 2018 by mehuljain